

CauSAIL: Causal Graph-Grounded Multimodal LLMs for Maritime Supply Chain Disruption Detection

Khan Enaet Hossain^{* †}
Department of Mathematics
York University
Toronto, Canada
enaet09@yorku.ca

Md. Sajeebul Islam Sk.[‡]
Department of CSE
BRAC University
Dhaka, Bangladesh
mdsajeebulislamsk@gmail.com

Md. Golam Rabiul Alam
Department of CSE
BRAC University
Dhaka, Bangladesh
rabiul.alam@bracu.ac.bd

Abstract

Maritime supply chain disruptions arise from complex interactions among vessel movements, port operations, satellite observations, public events, trade dependencies, corporate filings, and global pressure indicators. Existing monitoring systems often collapse these diverse modalities into static feature representations and rely on Retrieval-Augmented language models without sufficient causal grounding, modality control, and feature verification. We present CauSAIL that addresses these gaps through four connected modules: (1) integration of seven open-access datasets through a port-month alignment pipeline; (2) develop a Causal Temporal Supply Chain Knowledge Graph (CT-SCKG) and build Retrieval-Augmented Fine-Tuning (RAFT)-based feature sets to ground LLM decisions in structured and textual features; (3) develop a Suppression-Focal Modality-Aware LAFAM module (MA-LAFAM) to combine diverse modalities by handling missing modalities and class imbalance; (4) propose mathematical bias-mitigation methods to reduce seven IBM-recognised LLM bias types. Experimental results show that CauSAIL effectively identifies port-month disruption patterns by connecting vessel traces, satellite imagery, supply-chain relations, and retrieved documents. Ablation study demonstrates that vessel movement Automatic Identification System (AIS) features provide major disruption signal, and CT-SCKG, RAFT feature, MA-LAFAM, and mathematical bias-mitigation methods improve robustness and reduce unsupported LLM decisions. These experimental results suggest that reliable maritime intelligence systems benefit from grounded system design rather than large language model scale alone. We release our code here ¹.

^{*}Corresponding author

[†]These authors contributed equally to this work.

[‡]These authors contributed equally to this work.

¹<https://github.com/mdsajeebulislamsk/CauSAIL>

1 Introduction

Maritime supply chains are dynamic industrial systems. A disruption can arise from port congestion, vessel slowdowns, trade shocks, extreme weather, labor events, geopolitical conflict, supplier risk, or a combination of these factors (Shan and Schönberger, 2026). Automatic Identification System (AIS) tracks show vessel slowdown and dwell patterns (Shan and Schönberger, 2026). Satellite observations reflect changes near port waters (Avouris et al., 2025). Event streams report strikes, conflict, sanctions, and congestion (Shan and Schönberger, 2026). Trade records reveal dependencies among products and partner countries (Li et al., 2026). Company filings provide supplier and logistic risk (Shan and Schönberger, 2026).

Existing monitoring systems do not consider all the factors at the same time. Many reduce diverse features to static feature tables (Jiao et al., 2026) (Verschuur et al., 2025) (Shakil et al., 2025) and others depend on Retrieval-Augmented Language models with limited control over feature quality (Li et al., 2026) (Imran and Lee, 2026) (Liu et al., 2025). This existing workflow creates three research gaps. First, diverse features lack a shared decision across time and place. Second, the retrieved test includes irrelevant information, which raises the risk of unsupported LLMs outputs. Third, they do not address the real deployment settings because of missing modalities, severe label imbalance, and limited hardware budgets. We study these gaps and focus on five research questions.

- RQ1. How can diverse maritime supply chain feature be aligned into a single operational decision unit?
- RQ2. How can causal Knowledge Graphs (KG) and RAFT feature system reduce unsupported LLM decisions in logistics risk monitoring?

- RQ3. How can AIS vessel features, Sentinel-2 satellite features, KG embeddings, and LLM hidden states be integrated when some modalities are missing?
- RQ4. How can mathematical methods mitigate LLM bias types acknowledged by IBM?
- RQ5. Which data sources drive positive disruption detection, and which components reduce false positives?

To meet these demands, we propose Causal Supply-chain Artificial Intelligence for Logistics (CauSAIL). The main contributions of our study are as follows:

- We propose CauSAIL, a causal graph-grounded multimodal LLM framework for port-month maritime supply chain disruption detection.
- We build a port-month alignment pipeline that integrates seven open-access maritime and supply chain datasets into one operational decision unit.
- We develop a Causal Temporal Supply Chain Knowledge Graph (CT-SCKG) that connects ports, months, companies, products, trade flows, events, filings, and local observations in a time-aware structure to support grounded reasoning.
- We assemble RAFT-based feature sets with core and tangent documents that improve feature selection and reduce reliance on irrelevant text.
- We develop a Suppression-Focal Modality-Aware Fusion Module (MA-LAFAM) to combine AIS features, Sentinel-2 signals, KG embeddings, and LLM hidden states under missing modality and class imbalance settings.
- We propose mathematical bias-mitigation methods to reduce seven types of LLM bias and show that they improve robustness and reduce unsupported decisions.

2 Related Work

Prior work on maritime disruption detection was developed on individual modalities such as AIS movement (Martincic et al., 2021a), satellite imagery (Stepec et al., 2021a), public events (Huang

and Hampton, 2025a), trade records (Luo et al., 2025), company filings (Dolphin et al., 2026), and retrieved text (Ponnock et al., 2025).

2.1 Retrieval-Augmented Fine-Tuning (RAFT) Based Language Models

Retrieval-Augmented Generation (RAG) connects a parametric language model with an external document collection. Lewis et al. show that retrieval improves knowledge-intensive generation by giving the model access to non-parametric memory (Lewis et al., 2020). This idea fits supply-chain intelligence because relevant evidence changes across months, ports, firms, trade flows, and public events. Recent supply-chain studies use retrieval-augmented LLMs to connect external evidence with risk analysis, but retrieval noise remains a practical concern in operational settings (Li et al., 2026; Imran and Lee, 2026; Liu et al., 2025).

Retrieval-Augmented Fine-Tuning (RAFT) addresses this problem by training the model with relevant documents and distractor documents (Zhang et al., 2024). The relevant documents teach the model which evidence supports the answer. The distractor documents help model to ignore weak or unrelated context. CauSAIL follows this direction at the port-month level. It builds core evidence documents from AIS, Sentinel-2, KG, event, trade, and filing signals, and it builds tangent documents from weaker evidence. This design supports one practical goal: the LLM should learn which retrieved documents support the disruption decision, rather than trusting every retrieved passage.

2.2 Knowledge graph Grounding for LLM Decisions

Knowledge graphs (KGs) represent entities and relations as connected nodes and edges. Prior surveys show that KGs support entity linking, relation modeling, semantic search, and reasoning over structured knowledge (Hogan et al., 2021). Recent KG-LLM studies further show that graph structure helps language models ground their outputs in external knowledge instead of relying only on parametric memory (Pan et al., 2024). Prior supply chain studies use graphs to model firms, products, suppliers, and trade dependencies (Brintrup et al., 2020). These works show the value of relational structure for risk analysis. We use this method to link between vessels, ports, months, countries, companies, products, trade flows, filings, and events. A flat feature table stores these features in isolated

columns, but it does not preserve the relationships among them, and a text only LLM is also unstable to link among these entities. So, we develop a Causal Temporal Supply Chain Knowledge Graph (CT-SCKG) to connect Port, Month, Country, Company, Product, Trade flow, Event, Filing, Disruption Month, and Local Observation nodes. We use causal structure as domain grounding.

2.3 Multimodal Fusion for Maritime and Event Intelligence

Prior maritime studies use AIS traces to estimate vessel traffic, berth stay, congestion, port efficiency, and movement anomalies (Martincic et al., 2021b; Shakil et al., 2025). Satellite-based studies use optical or radar imagery to detect ships, observe port activity, and estimate maritime trade signals from space (Stepic et al., 2021b; Soldi et al., 2021; Jung et al., 2026). Event detection studies combine public reports, news streams, and structured event databases to identify conflict, sanctions, strikes, disasters, and logistics shocks (Leetaru and Schrodt, 2013; Kwak and An, 2014; Huang and Hampton, 2025b). These studies show the value of physical and textual evidence, but they work on each source separately. In CauSAIL, we address this gap through a modality-aware fusion design. We study AIS-derived vessel features, Sentinel-2 satellite indicators, CT-SCKG embeddings, and RAFT-guided LLM hidden states for a single industrial prediction task.

2.4 Bias-Aware Objectives and Reliability

LLM-based decision systems face bias from data, representation, retrieval, labels, and deployment context. IBM’s AI Fairness 360 work provides a framework for measuring and mitigating unwanted algorithmic bias in industrial settings (Bellamy et al., 2018). Recent LLM studies also report bias in generated content, evaluation behaviour, and decision support (Kotek et al., 2023; Stureborg et al., 2023). We address these risks through mathematical bias-mitigation methods. Orthogonal projection reduces direct reliance on port, time, or modality identity (Ali et al., 2026). MMD regularization reduces group-level representation shift (Gretton et al., 2012). Evidence alignment links LLM hidden states with KG evidence. Counterfactual consistency reduces unstable predictions under evidence changes. Calibration loss reduces overconfident outputs (Guo et al., 2017). Suppression-Focal loss focuses learning on hard disruption cases

and class imbalance (Lin et al., 2017). Fourth-Order Runge-Kutta (RK4) method hidden-state refinement stabilizes hidden representations before classification, following the use of differential equation style refinement in neural models (Chen et al., 2018).

2.5 Open LLM Backbones and Efficient Fine-Tuning

Open LLMs have made industrial NLP systems easier to deploy under local hardware and privacy constraints. Qwen2.5 provides instruction-tuned models across several sizes and reports strong performance in instruction following, structural data analysis, reasoning, and long-text generation (Yang et al., 2024). The Llama model provides open-weight language models for general language understanding and reasoning tasks (Dubey et al., 2024). Gemma 2 studies open models at practical sizes and focuses on efficient performance for downstream use (Gemma Team, 2024). Phi-3 studies compact language models for local and resource-constrained deployment (Abdin et al., 2024). GPT-OSS-20B provides a larger open-weight comparison point for local inference and private deployment (Bi et al., 2025). DeepSeek-VL provides an open vision-language model for real-world image-text understanding (Lu et al., 2024). GPT-2 remains a standard Transformer baseline for generative pretraining (Radford et al., 2019). Efficient adaptation methods reduce the cost of training these models. LoRA freezes the pretrained weights and trains low-rank adapter matrices, thereby reducing trainable parameters and memory usage (Hu et al., 2022). QLoRA extends this idea with 4-bit quantization, NF4 weights, double quantization, and paged optimizers, which makes large-model fine-tuning possible on smaller GPUs (Dettrmers et al., 2023). Hyperparameter search also matters in low-resource fine-tuning. Optuna provides a define-by-run optimization framework with pruning and sampling strategies for efficient hyperparameter selection (Akiba et al., 2019).

3 Methodology

We present the full workflow of CauSAIL in Appendix A.

3.1 Data Preprocessing, Feature Extraction, and Port-Month Alignment

We consider each port in each month as a single prediction sample. Let $p \in \mathcal{P}$ be a normalized

port key and let $t \in \mathcal{T}$ be a calendar month. We represent each sample as

$$x_{p,t} = \{X_{p,t}^{AIS}, X_{p,t}^{SAT}, X_t^{GSCPI}, X_{p,t}^{EVT}, X_{p,t}^{TRD}, X_{p,t}^{FIL}, X_{p,t}^{KG}, X_{p,t}^{TXT}\} \quad (1)$$

where $X_{p,t}^{AIS}$, $X_{p,t}^{SAT}$, X_t^{GSCPI} , $X_{p,t}^{EVT}$, $X_{p,t}^{TRD}$, $X_{p,t}^{FIL}$, $X_{p,t}^{KG}$, and $X_{p,t}^{TXT}$ denote AIS, Sentinel-2, global pressure, public event, trade, filing, graph, and retrieved-text evidence, respectively. The label is

$$y_{p,t} \in \{0, 1\} \quad (2)$$

where $y_{p,t} = 1$ denotes a disrupted port-month and $y_{p,t} = 0$ denotes a normal port-month.

By following this process, we match the temporal nature of maritime disruption. A port can be normal in one month and risky in another. The equation 1 also allows for lagged variables, chronological splits, and evidence selection within the same time window. We represent the seven datasets and their modalities in Table 1.

Before alignment, we extract modality-specific variables. AIS provides vessel observation count, unique vessel count, active days, mean speed, low-speed ratio, stationary ratio, and heading entropy. Sentinel-2 provides NDWI, NDVI, water fraction, vessel proxy, brightness, NIR, cloud cover, and valid pixel rate. GSCPI, GDELT, Comtrade, and SEC EDGAR provide pressure indicators, event signals, trade-dependency evidence, and company risk text for the construction of CT-SCKG and RAFT. We normalize port names to reduce duplicate identifiers. For example, New York/NJ maps to new york and new jersey, Norfolk maps to norfolk hampton roads, and Portland OR maps to portland oregon. We keep Los Angeles and Long Beach separate because their AIS and satellite patterns differ. We then map all records to the monthly level. AIS timestamps are aggregated into monthly vessel variables. Sentinel-2 acquisitions are assigned to their corresponding month. GSCPI is reported at the monthly level. GDELT events are grouped by event month. SEC filings use the filing month. Comtrade records are linked through country, product, and reporting period. Each record r_i is mapped by

$$\phi(r_i) \rightarrow (p, t) \quad (3)$$

where p is the normalized port key and t is the month. The final aligned dataset is

$$\mathcal{D} = \{(x_{p,t}, y_{p,t}) \mid p \in \mathcal{P}, t \in \mathcal{T}\} \quad (4)$$

where $\mathcal{D}$ denotes the full dataset, $x_{p,t}$ denotes the aligned representation for port p in month t , and $y_{p,t}$ denotes the disruption label. The sets $\mathcal{P}$ and $\mathcal{T}$ denote all normalized ports and month.

3.2 CT-SCKG and RAFT Evidence Construction

After alignment, we develop a Causal Temporal Supply Chain Knowledge Graph (CT-SCKG) to capture relations among ports, months, events, companies, products, trade flows, filings, and local observations. We define the graph as

$$\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathcal{R}, \tau) \quad (5)$$

where $\mathcal{V}$ is the node set, $\mathcal{E}$ is the edge set, $\mathcal{R}$ is the relation set, and τ stores the temporal context of each edge. Each edge is represented as

$$e_i = (v_s, r_i, v_o, t_i) \quad (6)$$

where v_s and v_o are source and target nodes, r_i is the relation type, and t_i is the associated month. For each port-month sample (p, t) , we retrieve the linked graph evidence $\mathcal{G}_{p,t}$ and encode it as a KG representation $h_{p,t}^{KG}$. This representation provides the model structured context about which events, filings, trade flows, products, and local observations are connected to the target port-month.

We also construct RAFT-based evidence for the LLM branch. For each query $q_{p,t}$, we build core documents $\mathcal{C}_{p,t}$ from relevant evidence and tangent documents $\mathcal{T}_{p,t}$ from weaker or distracting evidence. The RAFT input is

$$R_{p,t} = [q_{p,t}; \mathcal{C}_{p,t}; \mathcal{T}_{p,t}] \quad (7)$$

and the LLM produces the textual representation

$$h_{p,t}^{LLM} = f_{LLM}(R_{p,t}) \quad (8)$$

The equations 7 and 8 improve the model connecting structured supply-chain relations with retrieved textual evidence while reducing unsupported reliance on irrelevant documents (Zhang et al., 2024).

3.3 Modality Projection and LLM Encoding

After constructing CT-SCKG and RAFT evidence, we map all evidence sources into a shared hidden space. We use four main branches: AIS, satellite, KG, and RAFT-guided text. For each port-month sample (p, t) , the model receives

No.	Dataset	Modality	Source
1	GDELT Events 2.0	Event and text modality: conflict, trade, geopolitical events, sanctions, strikes, congestion, source URLs, and event tone.	(Leetaru and Schrodt, 2013)
2	UN Comtrade	Trade modality: bilateral trade flows, product dependency, partner-country exposure, and commodity-level risk signals.	(United Nations Statistics Division, 2024)
3	SEC EDGAR	Filing and text modality: 10-K and 8-K supply-chain risk mentions, supplier risk, production risk, and logistics risk.	(U.S. Securities and Exchange Commission, 2024)
4	NY Fed GSCPI	Economic-index modality: monthly global supply-chain pressure, z-score, lagged pressure, and rolling pressure context.	(Benigno et al., 2022)
5	U.S. DOT and NOAA AIS	Vessel-movement modality: dwell behaviour, speed, heading entropy, and local maritime anomaly signals.	(National Oceanic and Atmospheric Administration, 2024)
6	Sentinel-2 L2A optical imagery	Satellite modality: NDWI, NDVI, water fraction, vessel proxy, brightness, NIR, cloud cover, and valid pixel features.	(European Space Agency, 2015; Element 84, 2024)
7	NOAA MarineCadastre AIS	AIS modality: monthly vessel counts, active days, mean speed, low-speed ratio, stationary ratio, and heading entropy.	(NOAA Office for Coastal Management and Bureau of Ocean Energy Management, 2024)

Table 1: Seven datasets integrated into CauSAIL. Each dataset is mapped to the shared port-month key $k = (p, t)$.

$$\{X_{p,t}^{AIS}, X_{p,t}^{SAT}, h_{p,t}^{KG}, R_{p,t}\} \quad (9)$$

where $X_{p,t}^{AIS}$ and $X_{p,t}^{SAT}$ are structured physical features, $h_{p,t}^{KG}$ is the graph representation, and $R_{p,t}$ is the RAFT evidence input. Each non-text modality is projected into the LLM hidden dimension. We define

$$h_{p,t}^m = \text{Dropout}(\text{GELU}(W_m X_{p,t}^m + b_m)) \quad (10)$$

where $m \in \{AIS, SAT, KG\}$, W_m and b_m are modality-specific projection parameters, and $h_{p,t}^m$ is the projected hidden state. The RAFT evidence is encoded by the LLM as

$$h_{p,t}^{LLM} = f_{LLM}(R_{p,t}) \quad (11)$$

We use Qwen2.5 as the main language backbone and fine-tune it with QLoRA. The projected AIS, satellite, and KG states are then combined with the LLM hidden state:

$$H_{p,t} = \{h_{p,t}^{AIS}, h_{p,t}^{SAT}, h_{p,t}^{KG}, h_{p,t}^{LLM}\} \quad (12)$$

We apply gated cross-modal attention over $H_{p,t}$. It allows textual evidence to interact with physical and graphical evidence before final fusion. Now, the model connects vessel slowdown, satellite context, KG relations, and RAFT documents within the same port-month sample.

3.4 MA-LAFAM Fusion and Prediction

LaFAM is introduced as a label-free activation-map method for feature attribution in CNNs (Karjauv et al., 2024). We develop Suppression-Focal Modality-Aware LAFAM (MA-LAFAM) for port-month disruption prediction. MA-LAFAM combines modality experts, masking, gated fusion, and Suppression-Focal learning to handle missing evidence and class imbalance. For each port-month sample (p, t) , MA-LAFAM receives four hidden states:

$$H_{p,t} = \{h_{p,t}^{AIS}, h_{p,t}^{SAT}, h_{p,t}^{KG}, h_{p,t}^{LLM}\} \quad (13)$$

Each modality is first passed through a separate expert:

$$z_{p,t}^m = E_m(h_{p,t}^m); m \in \{AIS, SAT, KG, LLM\} \quad (14)$$

A gate network estimates the contribution of each modality:

$$g_{p,t} = \text{softmax}(W_g[h_{p,t}^{AIS}; h_{p,t}^{SAT}; h_{p,t}^{KG}; h_{p,t}^{LLM}]) \quad (15)$$

We use a mask $a_{p,t} \in \{0, 1\}^4$ to suppress missing evidence. The masked gate is

$$\tilde{g}_{p,t}^j = \frac{a_{p,t}^j g_{p,t}^j}{\sum_{\ell=1}^4 a_{p,t}^\ell g_{p,t}^\ell + \epsilon} \quad (16)$$

Algorithm 1 MA-LAFAM fusion

- 1: **Input:** $h^{AIS}, h^{SAT}, h^{KG}, h^{LLM}$ mask a
 - 2: **Output:** disruption probability $\hat{y}$
 - 3: $z^m \leftarrow E_m(h^m)$ for each modality m
 - 4: $g \leftarrow \text{softmax}(W_g[h^{AIS}; h^{SAT}; h^{KG}; h^{LLM}])$
 - 5: $\tilde{g} \leftarrow \text{Normalize}(a \odot g)$
 - 6: $h^* \leftarrow \sum_m \tilde{g}^m z^m$
 - 7: $\hat{y} \leftarrow \sigma(W_o h^* + b_o)$
 - 8: **return** $\hat{y}$
-

The fused representation and prediction are

$$h_{p,t}^* = \sum_{j=1}^4 \tilde{g}_{p,t}^j z_{p,t}^j \quad (17)$$

$$\hat{y}_{p,t} = \sigma(W_o h_{p,t}^* + b_o) \quad (18)$$

We train the module with Suppression-Focal loss:

$$\mathcal{L}_{SF} = -\alpha_y S(1 - p_y)^\gamma \log(p_y) \quad (19)$$

where p_y is the predicted probability of the true class, α_y weights rare disruption cases, γ focuses learning on difficult samples, and S suppresses majority-class samples. This process allows MA-LAFAM to weight the available modalities for each port-month and reduce the effect of weak evidence. We represent the fusion process in Algorithm 1.

3.5 Bias-Aware Hidden-State Regularization

We add bias-aware regularization to reduce shortcut dependence during training. In equation 20, $h_{p,t}$ denote the fused hidden representation for port p in month t . We first remove the component of $h_{p,t}$ that is aligned with a learned bias direction b :

$$h_{p,t}^\perp = h_{p,t} - \text{proj}_b(h_{p,t}) \quad (20)$$

The equation 20 is used to reduce direct dependence on identifiers such as port, time, or modality availability. We then train the model with the following objective:

$$\begin{aligned} \mathcal{L} = & \mathcal{L}_{SF} + \lambda_{mmd} \mathcal{L}_{MMD} + \lambda_{evid} \mathcal{L}_{evid} + \lambda_{cf} \mathcal{L}_{cf} \\ & + \lambda_{cal} \mathcal{L}_{cal} + 0.25 \mathcal{L}_{expert} + 0.10 \mathcal{L}_{gate} \end{aligned} \quad (21)$$

Here, $\mathcal{L}_{SF}$ handles class imbalance, $\mathcal{L}_{MMD}$ reduces distribution differences across groups, $\mathcal{L}_{evid}$ encourages agreement between LLM and KG representations, $\mathcal{L}_{cf}$ penalizes unstable predictions under small evidence changes, and $\mathcal{L}_{cal}$ reduces

overconfident probabilities. The expert and gate losses support modality-specific prediction and stable fusion. We provide details in Appendix D.

3.6 Reliability Checks

We evaluate reliability after training to test whether CauSAIL remains stable under uncertainty in evidence and model configuration. We use bootstrap confidence intervals, Brier score calibration, seed sensitivity, port-identity shortcut testing, time-leakage auditing, and tangent-document stress testing for evaluation. The detailed results are reported in Appendix C.

4 Experimental Setup

We present the experimental setup of CauSAIL in Table 2. We fine-tune Qwen2.5-3B-Instruct with QLoRA (Yang et al., 2024; Dettmers et al., 2023). We use Optuna 4.0 to select hyperparameter for LLM fine-tuning (Akiba et al., 2019). We report accuracy, precision, recall, F1, AUC, and MCC (Powers, 2011). Metric definitions are provided in Appendix E.

5 Results and Discussion

5.1 Main Test Performance

Table 3 reports the performance of CauSAIL on the test dataset, and it achieves 0.9062 accuracy, 0.8696 F1, 0.9610 AUC, and 0.7984 MCC. The recall score of 0.9091 indicates that CauSAIL captures most positive disruption cases. The precision of 0.8333 means that most predicted positive cases are correct. The MCC score of 0.7984 indicates a balanced classification result between disrupted and normal cases. The AUC value of 0.9610 shows that CauSAIL separates disrupted and normal port-months across decision thresholds.

Figure 1 shows the confusion matrix of CauSAIL. The model correctly identifies 19 normal port-months and 10 disrupted port-months. It misses one disruption case and raises two false disruption alerts. In an early-warning, false negatives are more costly because a missed disruption may delay response. False positives require to increase analysis of workload, but two false positives in 32 test cases are acceptable, as the system also provides KG and RAFT evidence for review.

5.2 Comparison with LLM Baselines

Table 4 compares CauSAIL performance with six LLM models. CauSAIL obtains the best accuracy,

Item	Setting	Item	Setting
Hardware	NVIDIA RTX 4070, 12 GB	Main backbone	Qwen2.5-3B-Instruct
Fine-tuning	QLoRA, 4-bit NF4, BF16	Hyperparameter search	Optuna 4.0
LoRA rank / alpha	16 / 32	Batch size	4
Learning rate	2×10^{-5}	Warmup / weight decay	0.20 / 0.02
Maximum epochs / patience	15 / 6	Dataset split	128 / 32 / 32
Primary metrics	Accuracy, precision, recall	Additional metrics	F1, AUC, MCC

Table 2: Training and evaluation setup for CauSAIL.

precision, F1, and AUC. Llama-3.2-3B and GPT-OSS-20B achieve perfect recall, but their lower precision shows that they produce more incorrect predictions of disruption. The comparison results show that larger LLMs do not automatically solve maritime disruption detection. GPT-OSS-20B is larger than the CauSAIL model, but it still falls behind CauSAIL in accuracy, precision, F1, and AUC. DeepSeek-VL achieves high recall but low precision and low accuracy, which means it tends to over-predict disruption. GPT-2 shows a different failure mode. Its AUC is higher than its F1, which suggests some ranking signal, but the thresholded decision is weak because recall is only 0.0909. CauSAIL improves the task by combining vessel movement evidence, satellite context, graph grounding, retrieval control, and modality-aware fusion rather than relying on language model scale alone.

5.3 Ablation Analysis

The ablation results of our study are presented in Table 5. The full model achieves the highest F1 and MCC. Removing AIS causes the largest performance collapse. In the NO_AIS setting, F1 drops from 0.8696 to 0.0000. This result shows that AIS carries the strongest positive disruption signal in the selected-local split. Vessel movement features such as dwell behaviour, low-speed ratio, stationary ratio, and heading entropy directly capture congestion port activity. From the ablation study, we get three main findings. First, AIS is necessary for positive disruption detection in this setting. AIS-only reaches 0.8000 F1, while NO_AIS, SAT_ONLY, and LLM_RAFT_ONLY fail to detect positive cases after thresholding. This means the local vessel movement signal is not only useful, but essential. Second, AIS alone is still weaker than the full model. AIS-only reaches 0.8000 F1 and 0.6880 MCC, while the full model reaches 0.8696 F1 and 0.7984 MCC. CT-SCKG links port-month states to events, filings, products, companies, and trade

flows. RAFT provides controlled textual evidence. MA-LAFAM learns how much weight each signal should receive. Third, the NO_RAFT_TEXT setting shows a precision-recall trade-off. Without RAFT text, precision increases to 1.0000, but recall drops to 0.6364. This means the model predicts fewer false positives, but it also misses more true disruptions. In contrast, the full model keeps a high recall while maintaining strong precision. Additional reliability checks are presented in Appendix C.




Figure 1: Confusion matrix of CauSAIL.

5.4 Discussion

The results support the main claim of the paper. Maritime disruption detection requires grounded multimodal system design. Local physical signals provide the strongest evidence. Structured relations and retrieved documents improve interpretation and reduce unsupported LLM responses. A larger LLM alone does not provide the same balance. CauSAIL is useful because each prediction can be connected back to port-month evidence. AIS explains movement-level disruption. Sentinel-2 adds visual and port-water context. CT-SCKG

Model	Accuracy	Precision	Recall	F1	AUC	MCC	TP	FP	FN	TN
CauSAIL	0.9062	0.8333	0.9091	0.8696	0.9610	0.7984	10	2	1	19

Table 3: Main selected-local test result for CauSAIL.

Model	Accuracy	Precision	Recall	F1	AUC
DeepSeek-VL	0.4062	0.3571	0.9091	0.5128	0.7359
GPT-OSS-20B	0.8125	0.6471	1.0000	0.7857	0.8918
Llama-3.2-3B	0.8438	0.6875	1.0000	0.8148	0.9394
Phi-3.5-mini	0.5312	0.3000	0.2727	0.2857	0.4545
Gemma-2B	0.5938	0.4500	0.8182	0.5806	0.6840
GPT-2	0.6562	0.5000	0.0909	0.1538	0.8268
CauSAIL (proposed)	0.9062	0.8333	0.9091	0.8696	0.9610

Table 4: Model performance comparison on the selected local port-month disruption detection task.

Modules					Performance							
SAT	AIS	KG	RAFT	LLM	Acc.	Prec.	Rec.	F1	AUC	AUPRC	MCC	$\Delta F1$
✓	✓	✓	✓	✓	0.9062	0.8333	0.9091	0.8696	0.9610	0.9399	0.7984	+0.0000
✓	✓	×	×	✓	0.8438	0.6875	1.0000	0.8148	0.9524	0.9063	0.7237	-0.0548
✓	✓	×	×	×	0.8438	0.6875	1.0000	0.8148	0.9524	0.9063	0.7237	-0.0548
×	✓	×	×	×	0.8438	0.7143	0.9091	0.8000	0.8961	0.7751	0.6880	-0.0696
✓	✓	✓	×	✓	0.8750	1.0000	0.6364	0.7778	0.9870	0.9751	0.7311	-0.0918
✓	×	✓	✓	✓	0.6562	0.0000	0.0000	0.0000	0.6970	0.4965	0.0000	-0.8696
✓	×	×	×	×	0.6562	0.0000	0.0000	0.0000	0.5844	0.4621	0.0000	-0.8696
×	×	×	✓	✓	0.6562	0.0000	0.0000	0.0000	0.6017	0.4553	0.0000	-0.8696

Table 5: Performance metrics for ablation studies with different CauSAIL module combinations.

links local behaviour to supply-chain entities and events. RAFT evidence helps the LLM distinguish core documents from tangent documents. The bias-aware and modality-aware components reduce over-confident reliance on one source. The confusion matrix provides that CauSAIL misses one disrupted port-month and creates two false positives. For an industrial monitoring team, this system can serve as a screening layer. Analysts can review the small number of flagged cases with attached KG and RAFT evidence. We provide an additional qualitative analysis in Appendix G to illustrate how correct and incorrect predictions are distributed within the learned test-space structure.

6 Limitations

CauSAIL still has several limitations for future work.

- Real-time deployment should be evaluated
- Computational resources should be expanded
- Short disruption windows should be tested

- Causal effects should be further tested

7 Conclusion

Our experiments show that maritime disruption detection improves when vessel movement, satellite signals, structured supply-chain relations, and retrieved textual evidence are integrated into a single port-month system. CauSAIL aligns seven data sources, grounds decisions with CT-SCKG and RAFT evidence, and fuses AIS, Sentinel-2, KG, and LLM states through MA-LAFAM. Results show stronger decision balance than LLM backbones alone, while ablations confirm that AIS is the strongest positive disruption signal and that graph grounding, retrieval control, and bias-aware fusion reduce unsupported decisions. Future work will scale CauSAIL to broader ports, real-time evidence streams, and analyst-in-the-loop validation.

References

Marah Abdin, Sam Ade Jacobs, Ammar Ahmad Awan, Jyoti Aneja, Ahmed Awadallah, Hany Awadalla,

- Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, Alon Benhaim, Mikhail Bilenko, Johan Bjorck, Sébastien Bubeck, and 1 others. 2024. [Phi-3 technical report: A highly capable language model locally on your phone](#).
- Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. 2019. [Optuna: A next-generation hyperparameter optimization framework](#). In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*.
- Z. Ali, T. Hussain, C.-L. Su, G. Parise, K. Sayler, M. Sadiq, A. D. Jurcut, and M. Elsis. 2026. [A novel hybrid signal processing based deep learning method for cyber-physical resilient harbor integrated ship-board microgrids](#). *IEEE Transactions on Industry Applications*.
- D. M. Avouris, F. Maciel, S. L. Sharp, S. E. Craig, A. G. Dekker, C. A. Di Vittorio, J. R. Gardner, E. Goldsmith, J. I. Gossn, S. R. Greb, B. K. Grunert, D. Gurlin, M. Jampani, R. M. Khan, B. Lowin, L. Mckinna, C. B. Mouw, I. Ogashawara, S. R. Calle, and J. Werdell. 2025. [Advancements in satellite observations of inland and coastal waters: Building towards a global validation network](#). *Remote Sensing*.
- Rachel K. E. Bellamy, Kuntal Dey, Michael Hind, Samuel C. Hoffman, Stephanie Houde, Kalapriya Kannan, Pranay Lohia, Jacquelyn Martino, Sameep Mehta, Aleksandra Mojsilović, Seema Nagar, Karthikeyan Natesan Ramamurthy, John Richards, Diptikalyan Saha, Prasanna Sattigeri, Moninder Singh, Kush R. Varshney, and Yunfeng Zhang. 2018. [Ai fairness 360: An extensible toolkit for detecting, understanding, and mitigating unwanted algorithmic bias](#). *arXiv preprint arXiv:1810.01943*.
- Gianluca Benigno, Julian di Giovanni, Jan J. J. Groen, and Adam I. Noble. 2022. [The gscpi: A new barometer of global supply chain pressures](#). Staff report, Federal Reserve Bank of New York.
- Xiao Bi, Deli Chen, Guanting Chen, Shanhuang Chen, Damai Dai, Chengqi Deng, Honghui Ding, Kai Dong, Qiushi Du, Zhe Fu, Huazuo Gao, Kaige Gao, Wenjun Gao, Ruiqi Ge, Kang Guan, Daya Guo, Jianzhong Guo, Guangbo Hao, Zhe Hao, and 2 others. 2025. [Gpt-oss-20b: Open-weight language model](#).
- Alexandra Brintrup, Yufei Wang, and Ashutosh Tiwari. 2020. [Supply networks as complex systems: A network-science-based characterization](#). *IEEE Systems Journal*.
- Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt, and David K. Duvenaud. 2018. [Neural ordinary differential equations](#). In *Advances in Neural Information Processing Systems*.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. [Qlora: Efficient finetuning of quantized llms](#). *Advances in Neural Information Processing Systems*.
- Rian Dolphin, Joe Dursun, Jarrett Blankenship, Katie Adams, and Quinton Pike. 2026. [Taxonomy-aligned risk extraction from 10-k filings with autonomous improvement using llms](#).
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, and 2 others. 2024. [The llama 3 herd of models](#).
- Element 84. 2024. [Earth search: STAC api for public earth observation data](#). Online catalog.
- European Space Agency. 2015. [Sentinel-2 user handbook](#). Technical report.
- Gemma Team. 2024. [Gemma 2: Improving open language models at a practical size](#).
- Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. 2012. [A kernel two-sample test](#). *Journal of Machine Learning Research*.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. [On calibration of modern neural networks](#). In *Proceedings of the 34th International Conference on Machine Learning*.
- Aidan Hogan, Eva Blomqvist, Michael Cochez, Claudia d'Amato, Gerard de Melo, Claudio Gutierrez, Sabrina Kirrane, José Emilio Labra Gayo, Roberto Navigli, Sebastian Neumaier, Axel-Cyrille Ngonga Ngomo, Axel Polleres, Sabbir M. Rashid, Anisa Rula, Lukas Schmelzeisen, Juan Sequeda, Steffen Staab, and Antoine Zimmermann. 2021. [Knowledge graphs](#). *ACM Computing Surveys*.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. [Lora: Low-rank adaptation of large language models](#). In *International Conference on Learning Representations*.
- Hsin-Hsiung Huang and Hayden Hampton. 2025a. [Forecasting geopolitical events with a sparse temporal fusion transformer and gaussian process hybrid: A case study in middle eastern and u.s. conflict dynamics](#).
- Hsin-Hsiung Huang and Hayden Hampton. 2025b. [Forecasting geopolitical events with a sparse temporal fusion transformer and gaussian process hybrid: A case study in middle eastern and u.s. conflict dynamics](#).
- IBM. 2024. [What is ai bias?](#) Online article.
- Muhammad Imran and Yugyung Lee. 2026. [Multimodal vision-language models in medical imaging: A survey of retrieval, interpretability, and trust](#). *IEEE Access*.

- Hang Jiao, Jincheng Gong, Huanhuan Li, Jasmine Siu Lee Lam, Yaqing Shu, Jin Wang, and Zaili Yang. 2026. [Llm4stp: A large language model-driven multi-feature fusion method for ship trajectory prediction](#). *Transportation Research Part E: Logistics and Transportation Review*.
- J. Jung and 1 others. 2026. [Watching maritime trade from space with satellite imagery](#). *Remote Sensing of Environment*.
- Adilet Karjaub and 1 others. 2024. [Lafam: Label-free activation maps for feature attribution](#).
- Hadas Kotek, Rikker Dockum, and David Q. Sun. 2023. [Gender bias and stereotypes in large language models](#). In *Proceedings of the ACM Collective Intelligence Conference*.
- Haewoon Kwak and Jisun An. 2014. [Understanding news geography and major determinants of global news coverage of disasters](#). In *Proceedings of the International AAAI Conference on Web and Social Media*.
- Kalev Leetaru and Philip A. Schrodt. 2013. [Gdelt: Global data on events, location, and tone, 1979-2012](#). *ISA Annual Convention*.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. [Retrieval-augmented generation for knowledge-intensive nlp tasks](#). In *Advances in Neural Information Processing Systems*.
- Huanhuan Li, Hang Jiao, Zhong Shuo Chen, Jasmine Siu Lee Lam, and Zaili Yang. 2026. [Covid crisis-aware maritime risk assessment: A bayesian network analysis](#). *Reliability Engineering & System Safety*.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. 2017. [Focal loss for dense object detection](#). In *Proceedings of the IEEE International Conference on Computer Vision*.
- Boyang Liu, Weiqi Li, Runjie Zhao, R. Ryan Klauer, Alan Hansen, Nicole Miller, Yixin Chen, Mark A. Blenner, and Yinjie J. Tang. 2025. [Leveraging retrieval-augmented generation to accelerate discoveries on mealworm larvae and plastic degradation](#). *Environmental Science & Technology*.
- Haoyu Lu, Wen Liu, Bo Zhang, Bingxuan Wang, Kai Dong, Bo Liu, Jingxiang Sun, Tongzheng Ren, Zhuoshu Li, Hao Yang, Yaofeng Sun, Chengqi Deng, Hanwei Xu, Zhenda Xie, Chong Ruan, and Zhanhui Dai. 2024. [Deepseek-vl: Towards real-world vision-language understanding](#).
- Wei Luo, Siyuan Kang, and Qian Di. 2025. [Global supply chain reallocation and shift under triple crises: A u.s.-china perspective](#).
- Tomaz Martincic, Dejan Stepec, João Pita Costa, Kristijan Cagran, and Athanasios Chaldeakis. 2021a. [Vessel and port efficiency metrics through validated AIS data](#). *CoRR*.
- Tomaz Martincic, Dejan Stepec, João Pita Costa, Kristijan Cagran, and Athanasios Chaldeakis. 2021b. [Vessel and port efficiency metrics through validated ais data](#). *arXiv preprint arXiv:2105.00063*.
- National Oceanic and Atmospheric Administration. 2024. [AIS data](#). Online dataset.
- NOAA Office for Coastal Management and Bureau of Ocean Energy Management. 2024. [MarineCadastre.gov: Automatic identification system vessel traffic data](#). Online dataset.
- Shirui Pan, Linhao Luo, Yufei Wang, Chen Chen, Jipapu Wang, and Xindong Wu. 2024. [Unifying large language models and knowledge graphs: A roadmap](#). *IEEE Transactions on Knowledge and Data Engineering*.
- Jesse Ponnock, Grace Kenneally, Michael Robert Briggs, Elinor Yeo, Tyrone Patterson III, Nicholas Kinberg, Matthew Kalinowski, and David Hechtman. 2025. [Real-time rag for the identification of supply chain vulnerabilities](#).
- David M. W. Powers. 2011. [Evaluation: From precision, recall and f-measure to roc, informedness, markedness and correlation](#). *Journal of Machine Learning Technologies*.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. [Language models are unsupervised multitask learners](#). *OpenAI Technical Report*.
- Md. Mostofa Nurannabi Shakil, Md. Alamgir Hossain, Md. Samiul Islam, and Nidadavolu Venkat Durga Sai Siva Vara Prasad Raju. 2025. [A novel deep learning approach for intrusion detection in maritime radar networks](#). *Scientific Reports*.
- Jing Shan and Jörn Schönberger. 2026. [Planning container flows through the eurasian rail network: Managing ad-hoc demand under limited capacity](#). *Omega*.
- Giovanni Soldi, Domenico Gaglione, Nicola Forti, Andrea Di Simone, Fabio C. Daffinà, Gianfranco Bottini, Danilo Quattrociochi, Leonardo M. Millefiori, Paolo Braca, Sandro Carniel, Peter Willett, Antonio Iodice, Daniele Riccio, and Giuseppe Ruello. 2021. [Space-based global maritime surveillance. part i: Satellite technologies](#). *IEEE Aerospace and Electronic Systems Magazine*.
- Dejan Stepec, Tomaz Martincic, and Danijel Skocaj. 2021a. [Automated system for ship detection from medium resolution satellite optical imagery](#). *CoRR*.

Dejan Stepec, Tomaz Martincic, and Danijel Skocaj. 2021b. [Automated system for ship detection from medium resolution satellite optical imagery](#). *arXiv preprint arXiv:2104.13923*.

Rickard Stureborg and 1 others. 2023. [Large language models are inconsistent and biased evaluators](#). *arXiv preprint*.

United Nations Statistics Division. 2024. [UN Comtrade Database](#). Online database.

U.S. Securities and Exchange Commission. 2024. [EDGAR: Electronic data gathering, analysis, and retrieval system](#). Online database.

Jasper Verschuur, J. Lumma, and Jim W. Hall. 2025. [Systemic impacts of disruptions at maritime chokepoints](#). *Nature Communications*.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jian Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, and 23 others. 2024. [Qwen2.5 technical report](#).

Tianjun Zhang and 1 others. 2024. [Raft: Adapting language model to domain specific rag](#).

A CauSAIL Workflow

We show the full CauSAIL workflow in Figure 2. It presents the complete path from data preprocessing to final port-month disruption prediction.

B Feature Extraction Details

For AIS, we aggregate vessel records by $k = (p, t)$. We compute observation count, unique vessel count, active days, mean speed over ground, low-speed ratio, stationary ratio, and heading entropy. The low-speed ratio is the fraction of AIS records below a speed threshold, and the stationary ratio is the fraction of records with near-zero movement.

For Sentinel-2, we summarize optical imagery over the port region. We compute NDWI and NDVI from spectral bands, then derive the water fraction, a vessel proxy, brightness statistics, NIR mean, cloud cover, and the valid pixel rate. These variables represent port-water and visible surface conditions.

For GSCPI, we use the current monthly value, z-score, absolute z-score, lag values, and rolling statistics. For GDELT, we keep disruption-relevant event records and summarize event pressure, mentions, tone, and event categories. For Comtrade, we create country-product trade dependency features. For SEC EDGAR, we extract filing text related to

supplier, production, logistics, and supply chain risk.

C Additional Reliability Simulations

We run additional reliability checks with a threshold fixed at 0.45. We examine metric uncertainty, probability calibration, seed sensitivity, port-identity shortcuts, time-leakage risk, and sensitivity to tangent RAFT documents.

Table 6 shows that the test split creates visible uncertainty around several metrics. Recall remains stable because the threshold model detects all positive cases in this robustness run. Precision, AUPRC, F1, and MCC show wider intervals, which means future evaluation should include larger and more diverse port-month samples.

Table 7 reports a Brier score of 0.1781. True positives and false positives have the same average predicted probability of 0.55, which suggests that the model separates many positive and negative cases but still assigns similar confidence to some incorrect positive predictions.

Table 8 shows sensitivity across two random seeds. The first seed gives stable disruption detection, while the second seed reduces recall and F1.

Table 9 reports no prediction flips after removing or shuffling port identity. This suggests that the model does not rely only on port labels for its thresholded predictions. The AUC decreases under the shuffled setting, so port identity may still affect ranking behaviour.

Table 10 shows that the model does not depend only on current-month pressure. The lagged-pressure setting improves several metrics in this run, while removing current pressure leaves F1 unchanged but lowers AUC. This suggests that temporal pressure variables affect ranking more than thresholded classification.

Table 11 shows that tangent-document changes do not flip predictions at the selected threshold. F1 and MCC stay fixed, while AUC and AUPRC vary. This means RAFT noise affects ranking confidence more than the final thresholded decision in this test.

D Bias-Aware Regularization Details

We map seven IBM-recognized LLM bias types to mathematical mitigation methods used in CauSAIL (IBM, 2024; Bellamy et al., 2018). The goal is to reduce shortcut learning, unstable evidence use, overconfidence, and class imbalance in port-month disruption prediction. The RAFT-based evidence




Figure 2: CauSAIL workflow for port-month maritime supply chain disruption detection.

Metric	Mean	Std.	95% CI Low	95% CI High
Accuracy	0.8429	0.0643	0.7188	0.9375
AUC	0.8809	0.0459	0.7857	0.9546
AUPRC	0.6842	0.1203	0.4000	0.8824
F1	0.8061	0.0894	0.5714	0.9375
MCC	0.7203	0.1000	0.5173	0.8819
Precision	0.6842	0.1203	0.4000	0.8824
Recall	1.0000	0.0000	1.0000	1.0000

Table 6: Bootstrap confidence intervals with $B = 1000$ resamples

Outcome	Count	Average predicted probability
TP	11	0.55
TN	16	0.35
FP	5	0.55
FN	0	–

Table 7: Brier score calibration summary. The Brier score is 0.1781.

terms are trained on relevant and distractor documents (Zhang et al., 2024). The regularization terms build on MMD (Gretton et al., 2012), calibration (Guo et al., 2017), focal loss (Lin et al., 2017), and neural ODE style refinement (Chen et al., 2018).

Orthogonal projection: Let h be a hidden representation and b be a bias basis learned from port, time, or modality indicators. We remove the component of h aligned with b :

$$\text{proj}_b(h) = \frac{h^\top b}{b^\top b + \epsilon} b, \quad (22)$$

$$h_\perp = h - \text{proj}_b(h). \quad (23)$$

The $h_\perp$ reduces direct reliance on shortcut variables.

RK4 hidden-state refinement: We refine hidden states with a learned neural dynamics function f_θ , following neural ODE representation refinement (Chen et al., 2018). The fourth-order Runge-Kutta update is

$$\begin{aligned} k_1 &= f_\theta(h_t), \\ k_2 &= f_\theta(h_t + \frac{dt}{2} k_1), \\ k_3 &= f_\theta(h_t + \frac{dt}{2} k_2), \\ k_4 &= f_\theta(h_t + dt k_3), \end{aligned} \quad (24)$$

$$h_{t+1} = h_t + \frac{dt}{6} (k_1 + 2k_2 + 2k_3 + k_4). \quad (25)$$

We use this step as representation refinement before classification.

MMD regularization: To reduce group-level representation shift, we apply Maximum Mean Discrepancy between two groups a and b (Gretton et al., 2012):

$$\mathcal{L}_{MMD} = \|\mu(h_a) - \mu(h_b)\|_2^2. \quad (26)$$

The groups can correspond to port partitions, time partitions, or modality-availability partitions.

Evidence alignment: We align LLM evidence with KG evidence using cosine distance:

$$\mathcal{L}_{evid} = 1 - \cos(h_{p,t}^{LLM}, h_{p,t}^{KG}). \quad (27)$$

The $\mathcal{L}_{evid}$ prevents the LLM branch from making decisions unsupported by graph evidence.

Counterfactual consistency: For a perturbed port-month sample $x_{p,t}^{cf}$, we minimize prediction divergence:

$$\mathcal{L}_{cf} = D_{KL}(P(y|x_{p,t}) \| P(y|x_{p,t}^{cf})). \quad (28)$$

This term reduces unstable decisions in response to small changes in evidence.

Calibration: We reduce overconfidence using a probability calibration loss, following prior work on neural network calibration (Guo et al., 2017):

$$\mathcal{L}_{cal} = \frac{1}{N} \sum_{i=1}^N (\hat{p}_i - y_i)^2, \quad (29)$$

Seed	Threshold	Accuracy	Precision	Recall	F1	AUC	AUPRC	MCC
1000	0.45	0.8438	0.6875	1.0000	0.8148	0.9394	0.8907	0.7237
1001	0.60	0.7500	0.5000	0.5000	0.5000	0.6667	0.4103	0.3333
Mean	–	0.7969	0.5938	0.7500	0.6574	0.8030	0.6505	0.5285
Std.	–	0.0663	0.1326	0.3536	0.2226	0.1928	0.3397	0.2761

Table 8: Seed sensitivity results

Setting	Accuracy	Precision	Recall	F1	AUC
Original	0.8438	0.6875	1.0000	0.8148	0.8810
Port ID removed	0.8438	0.6875	1.0000	0.8148	0.8810
Port ID shuffled	0.8438	0.6875	1.0000	0.8148	0.8420

Table 9: Port-identity shortcut test

where $\hat{p}_i$ is the predicted disruption probability and y_i is the true label.

Suppression-Focal loss: We use Suppression-Focal loss to handle hard examples and class imbalance. This follows the focal-loss idea of increasing attention to difficult and minority-class examples (Lin et al., 2017):

$$\mathcal{L}_{SF} = -\alpha_y S(1 - \hat{p}_y)^\gamma \log(\hat{p}_y), \quad (30)$$

where α_y gives higher weight to rare disruption cases, γ controls the focus on hard samples, and S suppresses easy majority-class examples.

Final objective: The full training objective is

$$\begin{aligned} \mathcal{L} = & \mathcal{L}_{SF} + \lambda_{mmd} \mathcal{L}_{MMD} + \lambda_{evid} \mathcal{L}_{evid} \\ & + \lambda_{cf} \mathcal{L}_{cf} + \lambda_{cal} \mathcal{L}_{cal} \\ & + 0.25 \mathcal{L}_{expert} + 0.10 \mathcal{L}_{gate}. \end{aligned} \quad (31)$$

The λ terms control the strength of each regularizer. This objective combines class-imbalance handling, evidence grounding, representation stability, calibration, and modality-gate regularization.

E Evaluation Metric Definitions

We report standard classification metrics. Let TP , TN , FP , and FN denote true positives, true negatives, false positives, and false negatives.

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}. \quad (32)$$

$$Precision = \frac{TP}{TP + FP}. \quad (33)$$

$$Recall = \frac{TP}{TP + FN}. \quad (34)$$

$$F1 = \frac{2 \cdot Precision \cdot Recall}{Precision + Recall}. \quad (35)$$

$$MCC = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP + FP)(TP + FN)} \cdot \sqrt{(TN + FP)(TN + FN)}}. \quad (36)$$

AUC measures ranking quality across decision thresholds. We use F1 and MCC as key metrics because disruption cases are fewer than normal cases.

F Full Ablation Results

Table 13 reports the complete ablation results, including classification scores, relative changes, and learned modality gate weights.

G Topological Visualization of Test Predictions

We include a topological visualization to inspect how the test samples are arranged in the learned representation space. Figure 3 projects the test port-month representations into a two-dimensional space. Each point represents one test sample. Edges connect nearby samples in the projected space, showing local neighbourhood structure. The colour of each point indicates the prediction outcome: true positive (TP), true negative (TN), false positive (FP), or false negative (FN). This figure is qualitative. We use it to understand the error pattern, not as a replacement for the quantitative scores reported in Table 3.




Figure 3: Topological space of maritime disruption states on the selected-local test split. Green points denote true positives, blue points denote true negatives, orange points denote false positives, and red points denote false negatives.

Setting	Accuracy	Precision	Recall	F1	AUC	AUPRC	MCC	$\Delta F1$	ΔAUC
Original	0.8438	0.6875	1.0000	0.8148	0.8810	0.6875	0.7237	0.0000	0.0000
Lagged pressure	0.9375	0.9091	0.9091	0.9091	0.9870	0.9778	0.8615	+0.0943	+0.1061
No current pressure	0.8438	0.6875	1.0000	0.8148	0.8268	0.6100	0.7237	0.0000	-0.0541

Table 10: Time-leakage audit

Tangent/core ratio	FP rate	F1	AUC	AUPRC	MCC
0	0.2381	0.8148	0.8810	0.6875	0.7237
1	0.2381	0.8148	0.8052	0.5991	0.7237
2	0.2381	0.8148	0.8874	0.6997	0.7237
3	0.2381	0.8148	0.8766	0.6799	0.7237

Table 11: Tangent-document stress test

The visualization shows three core findings. First, most test samples are correctly classified. The plot contains 10 true positives and 19 true negatives, which matches the confusion matrix result. This indicates that the learned representation separates many normal and disrupted port-months into stable local neighbourhoods.

Second, the disrupted cases with high predicted probabilities appear as clear positive-risk states. Several labelled true positives, such as Los Angeles, Miami, Baltimore, Tampa, Long Beach, and Norfolk, receive high disruption probabilities. These cases suggest that CauSAIL learns consistent risk patterns when vessel movement, satellite context, graph evidence, and RAFT evidence support the same decision.

Third, the error cases appear in boundary regions. The two false positives are close to nearby normal or mixed neighbourhoods, which suggests that the model identifies them as risk-like even though the ground truth label is normal. The single false negative appears away from the strongest positive-risk region. This pattern is important for deployment because it shows where the model may need additional evidence, better temporal context, or analyst review.

Overall, Figure 3 supports the main result analysis. CauSAIL not only produce strong aggregate scores. It also forms a representation space where most correct predictions occupy coherent local neighbourhoods, while errors are sparse and interpretable. This motivates the use of modality-aware fusion and reliability checks for port-month maritime disruption detection.

Bias type	Risk in CauSAIL	Mitigation method
Data bias	Some ports or months have richer AIS, satellite, or text evidence.	Modality masks, MMD regularization, and missing-modality checks.
Selection bias	Training samples may overrepresent specific ports or disruption periods.	Chronological split design and group-level regularization.
Representation bias	Hidden states may encode port, month, or modality identity shortcuts.	Orthogonal projection over port, time, and modality bias bases.
Confirmation bias	Retrieved text may reinforce weak or incorrect evidence of disruption.	RAFT tangent documents and KG-LLM evidence alignment.
Automation bias	The model may produce high-confidence unsupported alerts.	Calibration loss and evidence-linked prediction output.
Evaluation bias	A single split or metric may overstate reliability.	Ablation tests, bootstrap confidence intervals, and seed checks.
Deployment bias	Missing AIS, cloudy satellite scenes, or incomplete KG evidence may change behavior.	Modality masks and counterfactual consistency.

Table 12: Mapping between IBM-recognized LLM bias types and CauSAIL mitigation methods.

Table 13: Full ablation results on the test split. Δ values are measured relative to the full CauSAIL model. Gate values show the average modality contribution learned by MA-LAFAM.

Ablation	Acc.	F1	Prec.	Rec.	AUC	AUPRC	MCC	$\Delta F1$	$\Delta Acc.$	ΔAUC	g_{SAT}	g_{AIS}	g_{KG}	g_{LLM}
FULL_MODEL	0.9062	0.8696	0.8333	0.9091	0.9610	0.9399	0.7984	+0.0000	+0.0000	+0.0000	0.2677	0.1566	0.2829	0.2927
NO_KG_AND_RAFT	0.8438	0.8148	0.6875	1.0000	0.9524	0.9063	0.7237	-0.0548	-0.0625	-0.0087	0.3719	0.2062	0.0000	0.4218
LOCAL_ONLY_AIS_SAT	0.8438	0.8148	0.6875	1.0000	0.9524	0.9063	0.7237	-0.0548	-0.0625	-0.0087	0.3719	0.2062	0.0000	0.4218
AIS_ONLY	0.8438	0.8000	0.7143	0.9091	0.8961	0.7751	0.6880	-0.0696	-0.0625	-0.0649	0.0000	0.2193	0.0000	0.7807
NO_RAFT_TEXT	0.8750	0.7778	1.0000	0.6364	0.9870	0.9751	0.7311	-0.0918	-0.0312	+0.0260	0.2738	0.1529	0.2696	0.3037
NO_AIS	0.6562	0.0000	0.0000	0.0000	0.6970	0.4965	0.0000	-0.8696	-0.2500	-0.2641	0.3258	0.0000	0.3245	0.3497
SAT_ONLY	0.6562	0.0000	0.0000	0.0000	0.5844	0.4621	0.0000	-0.8696	-0.2500	-0.3766	0.4865	0.0000	0.0000	0.5135
LLM_RAFT_ONLY	0.6562	0.0000	0.0000	0.0000	0.6017	0.4553	0.0000	-0.8696	-0.2500	-0.3593	0.0000	0.0000	0.0000	1.0000